

COSC364 Assignment 1

George Drummond(53243258),
Ryan Cox(64656394)

May 4, 2017

Abstract

The following pertains to Assignment 1 of COSC364 S1 2017. This was a joint work by George Drummond and Ryan Cox in fulfilment of the course requirements of COSC364 S1 2017 and is a result of **equal contribution**.

Contents

1	Particularly well done aspects	1
2	Aspects to be improved	2
3	Atomicity of event processing	2
4	discussion	2
5	Main Program	5
6	RIPpacket	14
7	writelog	15

1 Particularly well done aspects

One of the initial successes of our design is the decent degree of modulation with which we were able to develop our program. Through the `RIP_packet.py` and `writelog.py` programs, we were able to achieve a reasonable level of abstraction around the tasks of packet processing and logging errors. This led to far easier reading and debugging.

Within these modulated programs, the methods of constructing and processing the RIP response packet was a strong point in our design. The packet was represented by a python string of hex values and Python's type conversion abilities were heavily utilised to allow for swift extraction of the relevant information. Again, this was able to be done under the hood by the helper program `RIP_packet.py` avoiding much unnecessary indexing by the main program.

These good practices led to a great degree of functionality being achieved by our program. Via the tiered testing method detailed below, we were able to quickly debug and achieve a working (and pretty!) result.

2 Aspects to be improved

One aspect of the program which could have been improved was the way in which we implemented our timing procedures. Though operationally sound, these are not as aesthetically pleasing as they could be, with the precise nature of their operation obscured to the casual observer. This implicit structure would perhaps have been better implemented by an explicit finite automata by defining state variables and functions and working with this as the basis for operation. This would perhaps not only have been nicer to look at but also easier to develop as we could have considered the protocol on a state by state basis rather than a conditional one.

3 Atomicity of event processing

Our program is atomic by nature in that each line is executed one at a time with nothing interrupting the program as it runs through the main round robin (fancy use of terminology) loop. In addition to this, the timers have a small element of randomness to them anyway and its not crucial that we react to them as soon as they tick over so we only react to them when we get back around to checking them. Every event has its time in the loop to do its thing and then waits while we get back around to it. Crucially, no event can interrupt another event that is running

(The blocking and waiting is the longest part of the time waiting of 0.5 seconds and the processing for the other stuff is about 0.0025 seconds so the chance we miss a packet is low, We are not blocking for %0.5 of the time (not doing select for %0.5 on average per program loop)(if we can even miss a packet and if we do we will get it back with the periodic update))

4 discussion

Our primary source for testing during development was our example topology of 7 routers configured as follows.




Figure 1: Sample Topology

Where metrics are shown in red and port numbers in blue. This network was implemented by the seven config files of the form `router x` $x \in \{1, 2, 3, 4, 5, 6, 7\}$ and, in its entirety, is large enough to allow for thorough testing of the RIP routing protocol. This example network is also of particular developmental use as it tests multiple configuration criteria, such as involving both the smallest and largest possible accepted port numbers, checking that the range worked correctly and wasn't accidentally off by one.

During early development, we routinely ran only routing demons 1 and 2 above. This enabled us to quickly check elements of basic functionality (sending/receiving packets, updating routing tables, timeout and garbage collection etc) but obviously did not give us much insight into the correctness of our protocol implementation.

Our second, deeper stage of testing (and the one used for the majority of development) involved a greater number of routing demons and multiple topological aspects. These were namely, the routers 1,2,3 and 6 above. This gave us a far more interesting, though still easily observable network of two transit routers 1 and 2 as well as two stub routers 3 and 6. By taking down either 3 or 6 we were able to test for the correctness of the timeout and garbage collection mechanisms as well as the propagation of link failure information in the form of triggered updates. Likewise, by taking down either 1 or 2, we were able to observe the isolation of a stub router (6 or 3 respectively) from the rest of the network. This stage of testing however, did stop short of the whole picture as it did not test for non-trivial path updates.

Finally a complete run of the network was used as our last stage of testing. Bringing up all routers allowed us to observe convergence times (MAYBE SOME DATA ON CONVERGENCE?) and also to check the validity of the shortest path information of each router. We could also, at this point, test the networks adaptation to extreme topological change. A favourite such change was to bring down routers 4 and 6. This had the joint effect of isolating 5 (i.e making router 5 unreachable from all other routers) and also changing a plethora of shortest path information (for example $D(7,3)$ changes from 10 to 11). It was with this particular test, that a great deal of valuable debugging was achieved.

A particularly useful tool for debugging was the `runlog_x` $x \in \{1, 2, 3, 4, 5, 6, 7\}$ files generated by the routing demon instance. Which after being initialised by the routing demon we could write the states and any errors in the routing demon so we could look back at what happened and when. Each routing demon makes its own separate log using its ID and since the log file object was part of the routing class it was easy to write to the log from the code without having to pass around the file object.

5 Main Program

```
1 #!/usr/bin/python
2 import sys
3 import select
4 import socket
5 import random
6 from RIP_packet import *
7 from writelog import *
8 import time
9
10 MAXBUFF = 600
11 MAXDATA = 512
12 INF = 16
13 HOSTID = '127.0.0.1'
14
15 # Router STATES: 0 -> Waiting for input with periodic updates
16 #                  1 -> Needs to send a triggered update
17
18 def valid_portn(portn):
19     return int(portn) in range(1024,64001)
20
21 def valid_ID(routerID):
22     return int(routerID) in range(1,64001)
23
24 def valid_metric(metric):
25     return int(metric) in range(0,INF+1)
26
27
28
29 class RIProuter:
30     def __init__(self, configFile):
31         self.periodic = 0
32         self.updateFlag = 0
33         self.configFile = configFile
34         self.parse_config()
35         self.socket_setup()
36         self.routingTable = RoutingTable(self.timers[1], self.timers[2]) #
37         timeout and garbage considered
38         self.log = init_log(self.routerID)
39
40         print('routerID =', self.routerID)
41         print('inport numbers =', self.inPort_numbers)
42
43         print('peerInfo =', self.peerInfo)
44         print('timers =', self.timers)
45         print('table=\n', self.routingTable)
46
47     def socket_setup(self):
48         self.inPorts = []
```

```

49         for portn in self.inPort_numbers:
50             newSocket = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
51
52             #newSocket.setsockopt(socket.SOL_SOCKET, socket.SO_REUSEADDR
, 1)
53
54             newSocket.bind((HOST_ID, portn))
55             self.inPorts += [newSocket]
56
57
58     def close_sockets(self):
59         ''' Close all sockets '''
60         for port in self.inPorts:
61             port.close()
62
63
64     def parse_config(self):
65         ''' Parse the supplied config file '''
66         lines = self.configFile.readlines()
67         for line in lines:
68             entries = line.split(',')
69             #print(entries)
70             lineType = entries[0]
71             tail = entries[1:]
72             if lineType == 'router-id':
73                 self.set_ID(tail)
74
75             elif lineType == 'input-ports':
76                 self.set_InPort_numbers(tail)
77
78             elif lineType == 'outputs':
79                 self.set_peerInfo(tail)
80
81             elif lineType == 'timers':
82                 self.set_timers(tail)
83
84
85     def set_ID(self, tail):
86         ''' Checks and stores routerID '''
87         myID = int(tail[0])
88         if valid_ID(myID):
89             self.routerID = int(tail[0])
90
91         else:
92             raise(IndexError('Router ID not valid'))
93
94
95     def set_InPort_numbers(self, tail):
96         ''' Checks and stores all supplied inport numbers '''
97         self.inPort_numbers = []
98         for portstring in tail:
99             port = int(portstring)

```

```

100         if (port not in self.inPort_numbers) and valid_portn(port):
101             self.inPort_numbers += [port]
102         else:
103             print("invalid inport port {} supplied".format(port))
104
105
106     def set_peerInfo(self, tail):
107         ''' Stores info relevent to immediate neighbours '''
108         self.peerInfo = dict()
109         for triplet in tail:
110             portN, metric, peerID = triplet.split('-')
111             if valid_portn(portN) and valid_metric(metric) and valid_ID(
peerID):
112                 self.peerInfo[int(peerID)] = (int(portN), int(metric))
113             else:
114                 print("invalid peer info for peer {}".format(peerID))
115
116
117     def set_timers(self, tail):
118         ''' Stores supplied timer info (i.e. periodic, timeout, garbage)
,,,
119         self.timers = []
120         for entry in tail:
121             self.timers += [int(entry)]
122
123
124
125     def send_updates(self):
126         ''' Sends an update message to each neighbour '''
127         i = 0
128         for peerID in self.peerInfo.keys():
129             print("update sent to {}".format(peerID))
130             write_to_log(self.log,
"Sent update to {}".format(peerID))
131             OutSock = self.inPorts[i] # use a different socket to send
each
132
133             peerPort = self.peerInfo[peerID][0]
134             response = self.response_packet(peerID)
135
136             OutSock.sendto(response.encode('UTF-8'), (HOST_ID, peerPort))
137
138             i += 1
139
140
141
142     def response_packet(self, peerID):
143         ''' Construct a response packet destined to a neighboring router.
Suitable for a periodic or triggered update '''
144
145         packet = ""
146         packet += rip_header(self.routerID)
147         for Entry in self.routingTable:

```

```

149         # Implement split horizon with poisson reverse
150         if (Entry.nextHop == peerID) or (Entry.garbageFlag == 1):
151             print("split horizon entry sent to {}".format(peerID))
152             packet += RTE(TableEntry(Entry.dest, INF, Entry.nextHop
)) # set metric to INF
153         else:
154             packet += RTE(Entry)
155
156     return packet
157
158 def process_rip_packet(self, packet):
159     ''' Processes a RIP response packet '''
160
161     (peerID, RTEs) = rip_packet_info(packet)
162
163
164     if not valid_ID(peerID):
165         print("[Error] peerID {} out of range".format(peerID))
166         write_to_log(self.log, "[Error] peerID {} out of range".
format(peerID))
167         #need to do something here
168
169     print("processing packet from {}".format(peerID))
170     cost = self.peerInfo[peerID][1]
171
172     # Consider direct link to peer Router
173     incomingEntry = self.routingTable.get_entry(peerID)
174     if incomingEntry is None:
175         print("added directlink entry to router {}".format(peerID))
176         self.routingTable.add_entry(peerID, cost, peerID)
177     else:
178         incomingEntry.metric = cost
179         incomingEntry.timeout = 0 # Reinitialise timeout for this
link
180         incomingEntry.garbageFlag = 0
181         incomingEntry.garbage = 0
182
183
184
185     for RTE in RTEs:
186         self.processRTE(RTE, peerID, cost)
187
188
189 def processRTE(self, RTE, peerID, cost):
190     '''processes an RTE of a RIP response packet from a peer router
,,,
191     (dest, metric) = RTE
192
193     new_metric = min(metric + cost, INF) # update metric
194
195     """check metric here?"""
196     currentEntry = self.routingTable.get_entry(dest)

```



```

197
198
199
200         if new_metric >= INF:
201             print("Path ({},{}) from {} not processed as unreachable".
format(dest, metric, peerID))
202             write_to_log(self.log,
203                 "Path ({},{}) from {} not processed as unreachable".
format(dest, metric, peerID))
204                 #do something here
205
206         if (currentEntry is None):
207             print("current route not in table")
208             if (new_metric < INF): # Add a new entry
209                 NewEntry = TableEntry(dest, new_metric, peerID)
210                 print('new Entry {}'.format(NewEntry))
211                 write_to_log(self.log,
212                     "New route added from {} to {} with Metric {}".
213                     .format(self.routerID, NewEntry, new_metric))
214
215
216                 self.routingTable.add_entry(dest, new_metric, peerID)
217
218         else: # Compare to existing entry
219
220             print("Existing entry for {}".format(dest))
221             if (currentEntry.nextHop == peerID): # Same router as
existing route
222
223                 currentEntry.timeout = 0 # Reinitialise timeout
224                 currentEntry.garbageFlag = 0
225                 currentEntry.garbage = 0
226
227                 if (new_metric != currentEntry.metric):
228                     self.existing_route_update(currentEntry,
new_metric, peerID)
229
230
231
232             elif (new_metric < currentEntry.metric):
233                 print("update route to {}".format(dest))
234                 write_to_log(self.log,
235                     "Route from {} to {} updated with new Metric
{}".
236                     .format(self.routerID, NewEntry, new_metric))
237                 self.existing_route_update(currentEntry, new_metric,
peerID)
238
239
240
241
242

```

```

243
244
245
246     def existing_route_update(self, currentEntry, new_metric, peerID):
247         ''' updates an existing routing table entry with a new metric '''
248         currentEntry.metric = new_metric
249         print("route to {} updated to metric = {}".format(currentEntry.
dest, new_metric))
250         currentEntry.nextHop = peerID
251
252         if (new_metric >= INF):
253             print("Triggered update flag set")
254             self.updateFlag = 1 #Set some update flag
255             currentEntry.garbageFlag = 1
256
257
258
259
260 class RoutingTable:
261     def __init__(self, timeoutMax, garbageMax):
262         self.table = []
263         self.timeoutMax = timeoutMax
264         self.garbageMax = garbageMax
265
266     def __iter__(self):
267         i = 0
268         while i < len(self.table):
269             yield(self.table[i])
270             i += 1
271
272     def __repr__(self):
273         blank = "-" * 54
274         print(blank + "\n| dest | metric | nextHop | flag | timeout |
garbage |")
275         for Entry in self.table:
276             print("{}{:>5} |{:>7} |{:>8} |{:>5} |{:>8.3f} |{:>8.3f} |".
format(
277                 Entry.dest, Entry.metric, Entry.nextHop, Entry.
garbageFlag,
278                 Entry.timeout, Entry.garbage))
279
280         return blank
281
282     def add_entry(self, dest, metric, nextHop):
283         self.table += [TableEntry(dest, metric, nextHop)]
284
285     def remove_entry(self, Entry):
286         print("Entry {} removed".format(Entry))
287         self.table.remove(Entry)
288
289     def get_entry(self, dest):
290         ''' returns required table entry if already present '''

```

```
291         i = 0
292         while i < len(self.table):
293             Entry = self.table[i]
294             if Entry.dest == dest:
295                 return Entry
296
297             i += 1
298
299         return None
300
301
302
303 class TableEntry:
304     def __init__(self, dest, metric, nextHop):
305         self.dest = dest
306         self.metric = metric
307         self.nextHop = nextHop
308         self.garbageFlag = 0
309         self.timeout = 0
310         self.garbage = 0
311
312     def __repr__(self):
313         return str((self.dest, self.metric,
314                     self.nextHop, self.garbageFlag, self.timeout, self.
315                     garbage))
316
317
318 def main():
319     configFile = open(sys.argv[1])
320     #configFile = open("router1.conf") # Just for developement
321     router = RIProuter(configFile)
322     selecttimeout = 0.5
323     periodicWaitTime = router.timers[0]
324
325     starttime = time.time() #Gets the start time before processing
326
327     while(1):
328         try:
329             # Wait for at least one of the sockets to be ready for
330             processing
331             print("table reads\n", router.routingTable)
332             readable, writable, exceptional = select.select(router.
333             inPorts, [], router.inPorts, selecttimeout) #block for incoming packets
334             for half a second
335
336             # Send triggered updates at this stage
337             if (router.updateFlag == 1):
338                 router.send_updates()
339                 router.updateFlag = 0
340
341             for sock in readable:
```

```

339         packet = sock.recv(MAX_BUFF).decode('UTF-8')
340         router.process_rip_packet(packet)
341
342     timeInc = (time.time() - starttime) #finds the time taken on
processing
343     #print("proc time = {}".format(timeInc))
344     starttime = time.time()
345     router.periodic += timeInc
346
347     if (router.periodic >= periodicWaitTime): # Periodic update
348         router.send_updates()
349         #Recalculate new random wait time in [0.8*periodic ,
1.2*periodic]
350         periodicWaitTime = random.uniform(0.8*router.timers
[0],1.2*router.timers[0])
351
352     router.periodic = 0 # Reset periodic timer
353     print("Periodic update")
354
355     for Entry in router.routingTable:
356
357         if (Entry.garbageFlag == 1):
358             Entry.garbage += timeInc
359             if (Entry.garbage >= router.timers[2]): # Garbage
collection
360                 print('Removed {}'.format(Entry))
361                 write_to_log(router.log,
362                             "[Warning] Route from {} to {}
has been removed"
363                             .format(router.routerID, Entry.
dest))
364                 router.routingTable.remove_entry(Entry)
365
366         else:
367             Entry.timeout += timeInc
368             if (Entry.timeout >= router.timers[1]): # timeout/
delete event
369                 print('Timeout')
370                 write_to_log(router.log,
371                             "[Warning] Route from {} to {}
has timed out"
372                             .format(router.routerID, Entry.
dest))
373                 Entry.metric = INF
374                 router.updateFlag = 1 # require triggered
update
375                 Entry.garbageFlag = 1 # Set garbage flag
376
377     except KeyboardInterrupt: # 'Taking down' router
378         print("Exiting program")
379         close_log(router.log)
380

```

```
381         router.close_sockets()
382         break
383
384
385 main()
```

6 RIPpacket

```

1  #!/usr/bin/python
2
3  # REMEMBER bytes.hex() and bytes.fromhex()
4
5  def bytes_to_int(byte_string):
6      return int.from_bytes(byte_string, byteorder='big')
7
8  #def int_to_bytes(myint, size):
9      #''' converts integer to 'size' number of bytes '''
10     #return (myint).to_bytes(size, byteorder='big').hex()
11
12
13 def int_to_bytes(myint, size):
14     suffix = hex(myint)[2:]
15     prefix = '0'*(2*size-len(suffix))
16     return (prefix + suffix)
17
18 def rip_header(routerID):
19     header = '0201' + int_to_bytes(routerID, 2)
20     return header
21
22 def RTE(Entry):
23     zero_row = int_to_bytes(0, 4)
24     s = zero_row
25     s += int_to_bytes(Entry.dest, 4)
26     s += zero_row
27     s += zero_row
28     s += int_to_bytes(Entry.metric, 4)
29     return(s)
30
31
32 def rip_packet_info(packet):
33     ''' Extracts relevent info from a RIP response packet '''
34
35     RTEs = []
36
37     peerID = int(packet[4:8], 16)
38
39     i = 8 # Start of first RTE
40     while i < len(packet):
41
42         dest = int(packet[i+8:i+16], 16) # Read dest from RTE
43         metric = int(packet[i+32:i+40], 16) # Read metric from RTE
44         RTEs += [(dest, metric)]
45
46         i += (8*5) # Proceed to next RTE
47
48     return(peerID, RTEs)

```

7 writelog

```
1 from time import *
2 import inspect
3
4 def init_log(ID):
5     """Initialises the log with an ID for the file"""
6     filename = "runlog_" + str(ID) + ".log"
7     program_log = open(filename, 'w')
8     program_log.write("Log File for {} in program {}\n{}\n"
9                       .format(ID, inspect.stack()[1][1], "/" * 200))
10    write_to_log(program_log, "Log Started")
11    return program_log
12
13 def write_to_log(log, string):
14     """takes a log object and writes the given string and timestamps it"""
15     logtime = strftime("[%H:%M:%S %d/%m/%Y] ", gmtime())
16     log.write(logtime + string + '\n')
17
18 def close_log(log):
19     """closes the file"""
20     write_to_log(log, "Log Ended")
21     log.close()
22
23 '''#Test case
24 log = init_log(0)
25 write_to_log(log, "Error 1")
26 write_to_log(log, "Error 2")
27 write_to_log(log, "Error 3")
28 write_to_log(log, "Error 4")
29 close_log(log)'''
```